



Figure 3. DFlash inference performance by dataset and checkpoint (hatched VDC bars = cross-GPU exploratory comparison)


MVBENCH — end-to-end latency



MVBENCH — throughput



VDC50 — end-to-end latency (*)



VDC50 — throughput (*)

