

Figure 5. Speculative inference time and end-to-end speedup (hatched VDC bars = cross-GPU exploratory comparison)

